

Epistemic Uncertainty in LLM Hallucinations

CENG463 – Technical Checkpoint Progress Report

May 11, 2026

Group Members

Ata TürkYılmaz	300201055
Berkhan Kargılı	290201108
Atakan Kumbaracı	300206012

<https://github.com/TAtaTrkylmz/CENG463-FinalProject>

1. Problem Definition

Task. This project investigates whether *epistemic uncertainty signals* extracted from large language models (LLMs) can reliably detect hallucinated outputs. Given a question and a candidate answer, the system must classify the answer as factual (label 0) or hallucinated (label 1). The task is therefore a **binary classification** problem framed around natural language understanding and model introspection.

Importance. Hallucinations – cases where a language model generates fluent but factually incorrect content – represent one of the most critical reliability problems in deployed NLP systems. Detecting them automatically is essential for safe deployment in high-stakes domains such as medicine, law, and education. Rather than relying on external fact-checking, this project explores whether the model’s own uncertainty (expressed via token-level log-probabilities) can serve as an intrinsic hallucination signal, a direction that is both computationally efficient and broadly applicable.

Dataset, Input, and Output. The project uses the **HaluEval QA** benchmark. Each original sample contains a question, a knowledge snippet, a factual answer, and a hallucinated answer. During preprocessing, each sample is expanded into two supervised rows, giving a balanced binary classification dataset. The system input is a (question, candidate answer) pair – optionally augmented with the knowledge snippet as context – and the output is a binary hallucination label.

Technical Challenges. The task is technically challenging for several reasons. First, hallucinated answers are often fluent and superficially plausible, making surface-level detection difficult. Second, uncertainty signals derived from small local models (e.g., `distilgpt2`) may not transfer reliably to the larger models that produce hallucinations in practice. Third, distinguishing genuine epistemic uncertainty from aleatoric uncertainty (inherent ambiguity in the question) is a non-trivial research problem. Finally, the binary construction of HaluEval introduces its own biases that must be understood during evaluation.

2. Dataset and Preprocessing Status

- **Dataset name and source:** HaluEval QA, available via the Hugging Face `datasets` library (`pminervini/HaluEval`, split `qa`).
- **Sample construction:** Each original HaluEval record yields two supervised examples – one factual (label 0) and one hallucinated (label 1) – resulting in a balanced binary dataset.
- **Features / input type:** Text fields: `question`, `answer` (candidate), and optionally `knowledge` (the reference snippet used in context-mode inference).
- **Target label:** Binary integer – 0 (factual) or 1 (hallucinated).
- **Train / validation / test split:** Produced deterministically by `scripts/prepare_dataset.py`. Splits are written as JSONL files under `data/processed/halueval_qa/`.
- **Preprocessing steps completed:**
 - Download and normalisation via Hugging Face `datasets`.
 - Expansion into two supervised rows per original sample.

- Deterministic train / val / test split and serialisation to JSONL.
- Processed splits verified to exist under `data/processed/halueval_qa/`.
- **Known issues:** The dataset is artificially balanced by construction, which may not reflect real-world hallucination rates. Downstream metrics should be interpreted with this in mind. No missing-value or encoding issues were observed.

Status: The dataset has been downloaded, normalised, split, and is ready for experiments.

3. Literature Review Progress

The literature review is ongoing, focusing heavily on recent advancements in hallucination detection and uncertainty quantification. The following thematic areas have been surveyed:

Hallucination Taxonomy and Detection Frameworks. Huang et al. [1] provide a comprehensive taxonomy of hallucinations in modern LLMs, distinguishing between factuality issues and faithfulness issues. Their survey confirms the urgency of moving beyond simple text-overlap metrics toward intrinsic model evaluation, which directly motivates our use of probability distributions as a detection signal.

Epistemic Uncertainty and Token Probabilities. Estimating predictive uncertainty in LLMs has recently shifted toward exploiting internal log-probabilities. Ling et al. [2] formally separate aleatoric and epistemic uncertainty during in-context learning, validating that token-level probability variations effectively signal a model’s lack of knowledge. We operationalize this finding in our entropy-based baseline. Furthermore, Das and Srihari [3] demonstrated that multi-faceted LLM ensembles utilizing role-based prompting can robustly detect these hallucinations, inspiring the architecture of our proposed hybrid method.

Real-World LLM Deployment and Reliability Risks. The necessity of our framework is underscored by recent attempts to deploy generative models in critical, domain-specific applications. For instance, Prado et al. [4] explore embedded LLM agents in aviation flight decks, Cao et al. [5] propose hybrid models for energy-safe industrial monitoring, and Abbas and Atwell [6] utilize LLMs for automated student assessment. Across diverse domains—from education to Arabic language fine-tuning [7]—the lack of reliable hallucination detection remains a primary bottleneck, justifying our focus on uncertainty-aware classification.

4. Baseline Models

Three baselines have been selected and fully implemented:

1. **Lexical SVM** (`lexical_svm`). A TF-IDF vectoriser over the concatenated question and answer, fed into a linear Support Vector Machine. This baseline is intentionally shallow: it captures surface-level lexical overlap that may correlate with hallucination

(e.g., vague or unusually specific phrasing) without any LM involvement. It serves as the interpretable lower bound.

2. **Entropy Classifier (entropy)**. Token-level log-probabilities are extracted from a local causal LM (`distilgpt2`) running in *memory mode* (no retrieved context). Aggregate features – mean log-probability, total sequence entropy, and token count – are passed to a logistic regression classifier. This baseline directly operationalises the epistemic-uncertainty hypothesis: a model that is uncertain about the tokens in a hallucinated answer should assign lower probabilities to those tokens.
3. **RAG Compare (rag_compare)**. The same local LM scores each answer twice – once without context (*memory mode*) and once with the HaluEval knowledge snippet prepended (*context mode*). The delta in mean log-probability between the two modes is thresholded to produce a binary prediction. A factual answer is expected to receive a larger probability boost from relevant context than a hallucinated one.

All three baselines are implemented in `src/llm_uncertainty/baselines.py` and are runnable via `scripts/run_baseline.py`. Preliminary results are reported in Section 5.

5. Initial Experimental Results

Initial experiments were conducted on a small subset (80 training / 20 validation samples) to verify the end-to-end pipeline. Results are summarised in Table 1.

Table 1: Preliminary validation results on HaluEval QA.

Model	Accuracy	Precision	Recall	F1	AUROC
Lexical SVM	0.943	0.944	0.943	0.943	0.972
Entropy	0.929	0.930	0.929	0.929	0.970
RAG Compare	0.751	0.751	0.751	0.751	0.791

Analysis of Initial Results: The preliminary metrics demonstrate that the pipeline is fully operational. The **Entropy Baseline** performs exceptionally well (AUROC 0.970), validating our core hypothesis that internal token probabilities serve as a strong signal for hallucination detection. Interestingly, the **Lexical SVM** achieves an unusually high accuracy (94.3%), which strongly suggests the presence of lexical artifacts or biases within the HaluEval dataset itself. The **RAG Compare** method establishes a solid, functioning baseline (AUROC 0.791) following threshold inversion, confirming that relevant context improves the model’s confidence for factual answers.

6. Planned Improvements and Technical Direction

The next phase will proceed in three directions:

Scale up experiments. Run all three baselines on the full HaluEval QA dataset (removing the `-limit` flag) to obtain stable, reportable metrics.

Stronger LM backbone. Replace `distilgpt2` with a larger instruction-tuned model (e.g., GPT-2 medium or an open-weight instruction model) to obtain better-calibrated log-probabilities for the entropy and RAG-compare baselines. This directly tests whether epistemic uncertainty scales with model capacity.

Expanded uncertainty features. Beyond mean log-probability and sequence entropy, add token-level variance, top- k probability mass, and calibration metrics (Expected Calibration Error, Brier score) as features for the entropy classifier.

Threshold optimisation for RAG Compare. The current delta-thresholding approach uses a fixed threshold. A calibrated threshold search on the validation set will be introduced to improve precision-recall balance.

Addressing potential class imbalance. Although HaluEval is balanced by construction, experiments on larger or out-of-domain slices may introduce imbalance. Oversampling (SMOTE) or class-weighted loss will be evaluated if needed.

Implementation of the Proposed Hybrid Method. To satisfy the core project requirement of proposing a novel method, the next immediate technical step is the development of a **Hybrid Uncertainty-Aware SVM**. This model will concatenate semantic features (textual TF-IDF embeddings or RAG similarity scores) with epistemic features (Shannon entropy, token-level variance, and log-probabilities). By combining surface-level representations with deep model introspection, we hypothesize this hybrid approach will outperform individual baselines and better handle datasets lacking lexical biases.

7. Ablation and Error Analysis Plan

Ablation study. The following components will be ablated in the final report:

- *Context vs. memory prompting:* Removing the knowledge snippet from the RAG-compare pipeline to measure the contribution of retrieved context.
- *Feature subsets for entropy classifier:* Evaluating subsets of log-probability features (mean only vs. mean + variance vs. full feature set) to understand which uncertainty dimensions are most predictive.
- *LM backbone size:* Comparing `distilgpt2` against larger models to quantify the effect of model capacity on uncertainty signal quality.

Error analysis. Model failures will be examined through:

- *Confusion matrices:* To identify asymmetric error patterns (false positives vs. false negatives).
- *Qualitative inspection:* Manual review of misclassified examples to identify recurring failure modes (e.g., hallucinations that are nearly paraphrases of factual answers).
- *Confidence-error correlation:* Cases where the model assigns high confidence to a wrong prediction will be flagged as overconfident failures and examined separately.
- *Difficulty stratification:* Errors will be stratified by question type and answer length to identify subgroup weaknesses.

8. Visualization and Interpretation

The reporting pipeline successfully generates comprehensive visual diagnostics. The following visualisations are automatically generated and stored in `reports/figures/`:

- **Confusion matrices** exposing false-positive and false-negative patterns for all baselines.
- **ROC and Precision-Recall curves** comparing the baselines under varying decision thresholds.
- **Calibration plots** (reliability diagrams) assessing the alignment between predicted probabilities and empirical accuracy.

As an example of the pipeline output, Figure 1 shows the generated ROC curves for our initial baselines.




Figure 1: Validation ROC curves generated by the reporting pipeline for the Lexical SVM, Entropy Classifier, and RAG Compare baselines.

9. GitHub and Reproducibility Status

- **Repository:** <https://github.com/TataTrkylmz/CENG463-FinalProject>
- **README:** Present and describes project goals, default dataset, setup, and a complete step-by-step end-to-end reproduction workflow.
- **Dependency file:** `requirements.txt` is present and pins all major dependencies: `pandas`, `numpy`, `scikit-learn`, `pyyaml`, `tqdm`, `matplotlib`, `seaborn`, `datasets`, `evaluate`, `transformers`, `torch`.

Reproduction instructions:

1. `pip install -r requirements.txt`
2. `python scripts/prepare_dataset.py -limit 100`
3. `python scripts/run_inference.py -input data/processed/halueval_qa/train.jsonl \`
`-output results/generations/halueval_qa_train_memory.jsonl -mode memory -limit 80`
4. `python scripts/run_baseline.py -baseline entropy ...`
5. `python scripts/make_report_assets.py`

10. Current Challenges and Next Steps

Current challenges:

- *LM scoring speed:* Extracting token-level log-probabilities for every sample with a local causal LM is slow on CPU. Running the full dataset currently requires either GPU access or significant wall-clock time.
- *Model quality vs. availability:* `distilgpt2` is fast and freely available but is not instruction-tuned, meaning its log-probabilities may not faithfully reflect semantic hallucination. A stronger model requires more compute.
- *Metric consolidation:* Small-sample runs have been used only for pipeline validation; final numeric results are not yet stable.
- *Literature review depth:* The `docs/papers/` folder is a placeholder; written summaries of reviewed papers remain to be committed.

Next steps before final submission:

1. Run all baselines on the full HaluEval QA dataset and record stable metrics.
2. Evaluate a larger or instruction-tuned LM backbone for the entropy and RAG-compare baselines.
3. Expand uncertainty features (ECE, Brier score, token-level variance).
4. Complete ablation experiments and document findings.
5. Perform qualitative error analysis and produce all planned visualisations.
6. Commit written literature review summaries to `docs/papers/`.
7. Increase commit frequency to maintain a consistent development record.

10. References

- [1] Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., ... & Liu, T. (2025). *A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions*. ACM Computing Surveys.
- [2] Ling, C., Zhao, X., Zhang, X., Cheng, W., Liu, Y., Sun, Y., ... & Chen, H. (2024). *Uncertainty Quantification for In-Context Learning of Large Language Models*. arXiv preprint arXiv:2402.10189.
- [3] Das, S., & Srihari, R. K. (2024). *Compos Mentis at SemEval2024 Task6: A Multi-Faceted Role-based Large Language Model Ensemble to Detect Hallucination*. Proceedings of the 18th International Workshop on Semantic Evaluation (SemEval-2024).
- [4] Prado, J. P. D. A., Mangion, D. Z., Muscat, A., Mizzi, S., Manicolo, A., Zammit, B., & Gauci, J. (2025). *Enhancing Flight Deck Decision Support with Distributed GenAI: A Multi-Agent Approach*. 2025 IEEE Aerospace Conference.
- [5] Cao, Q., Li, C., & Shi, H. (2025). *Toward Energy-Safe Industrial Monitoring: A Hybrid Language Model Framework for Video Captioning*. Applied Sciences, 15, 12848.
- [6] Abbas, N., & Atwell, E. (2025). *Cognitive Computing with Large Language Models for Student Assessment Feedback*. Big Data and Cognitive Computing, 9, 112.
- [7] Khader, K. A., Hussein, M. S., & Abu-Issa, A. S. (2025). *Adapting Large Language Models for Arabic: Comparative Evaluation, Fine-Tuning, and Ethical Deployment*. IEEE Access, 13, 182632-182650.